

1. General Usage

Open the terminal under the RankLib directory and type in

```
> java -jar bin/RankLib.jar
```

You should see all necessary parameters as below Parameters in the square brackets are optional with their default value shown at the end of the line

Usage java -jar RankLib.jar [Params]	
Params	
[-h] Training tuning and evaluation	
-t train file	Training data
-a ranker type	Specify which ranking algorithm to use
	0 MART gradient boosted regression tree
	1 RankNet
	2 RankBoost
	3 AdaRank
	4 Coordinate Ascent
	6 LambdaMART
	7 ListNet
	8 Random Forests
-f feature file	Feature description file list features to be considered by the learner each on a separate line If not specified all features will be used
-m metric2t metric	Metric to optimize on the training data Supported MAP NDCG k DCG k P k ERR k ERR k default ERR 10
-g gmax label	Highest judged relevance label It affects the calculation of ERR default 4 i.e. 5 point scale 0 1 2 3 4
-s [-silent]	Do not print progress messages which are printed by default
-v validate file	Specify if you want to tune your system on the validation data default unspecified If specified the final model will be the one that performs best on the validation data
-tvs tvs x in 0 1	Set train validation split to be x 1 0 x
-s save model	Save the learned model to the specified file default not save
-f test file	Specify if you want to evaluate the trained model on this data default unspecified
-tts tts x in 0 1	Set train test split to be x 1 0 x tts will override tvs
-m metric2T metric	Metric to evaluate on the test data default to the same as specified for metric2t
-n norm method	Normalize feature vectors default no normalization Method can be sum normalize each feature by the sum of all its values

~~/score~~ normalize each feature by its mean standard deviation

~~/linear~~ normalize each feature by its min max values

~~/kcv~~ k

Specify if you want to perform k fold cross validation using ONLY the specified training data default NoCV

~~-tvs~~ can be used to further reserve a portion of the training data in each fold for validation

~~/kcvmd~~ dir

~~Directory~~ for models trained via cross validation default not save

~~/kcvmn~~ model

Name for model learned in each fold It will be prefix ed with the fold number default empty

- [-] RankNet specific parameters

~~/epoch~~ T

~~The~~ number of epochs to train default 100

~~/layer~~ layer

~~The~~ number of hidden layers default 1

~~/node~~ node

~~The~~ number of hidden nodes per layer default 10

~~/lr~~ rate

~~Learning~~ rate default 0.00005

- [-] RankBoost specific parameters

~~/round~~ T

~~The~~ number of rounds to train default 300

~~/c~~ k

~~Number~~ of threshold candidates to search 1 to use all feature values default 10

- [-] AdaRank specific parameters

~~/round~~ T

~~The~~ number of rounds to train default 500

~~[]-noeq~~

~~Train~~ without enqueueing too strong features default unspecified

~~/tolerance~~ t

~~Tolerance~~ between two consecutive rounds of learning default 0.002

~~/max~~ times

The maximum number of times can a feature be consecutively selected without changing performance default 5

- [-] Coordinate Ascent specific parameters

~~/r~~ k

~~The~~ number of random restarts default 5

~~/iteration~~

~~The~~ number of iterations to search in each dimension default 25

~~/tolerance~~ t

~~Performance~~ tolerance between two solutions default 0.001

~~/reg~~ slack

~~Regularization~~ parameter default no regularization

}- [-] {MART LambdaMART specific parameters

~~/tree~~ t

~~Number~~ of trees default 1000

~~/leaf~~ l

~~Number~~ of leaves for each tree default 10

~~/shrinkage~~ factor

~~Shrinkage~~ or learning rate default 0.1

~~/c~~ k

Number of threshold candidates for tree splitting 1 to use all feature values default 256

~~/m~~ls n

~~Min~~ leaf support minimum samples each leaf has to contain default 1

~~/estop~~ e

Stop early when no improvement is observed on validation data in e consecutive rounds default 100

- [-] Random Forests specific parameters	
bag r	N umber of bags default 300
rate r	S ub sampling rate default 1 0
rate r	F eature sampling rate default 0 3
type type	R anker to bag default 0 i e MART
tree t	N umber of trees in each bag default 1
leaf l	N umber of leaves for each tree default 100
shrinkage factor	S hrinkage or learning rate default 0 1
c k	Number of threshold candidates for tree splitting 1 to use all feature values d efault 256
min s n	M in leaf support minimum samples each leaf has to contain default 1
stop e	Stop early when no improvement is observed on validation data in e consecutive r ounds default 100
[+] Testing previously saved models	
load model	The model to load
test file	T est data to evaluate the model specify either this or rank but not both
rank file	R ank the samples in the specified file specify either this or test but not both
metric2T metric	M etric to evaluate on the test data default ERR 10
gmax label	H ighest judged relevance label It affects the calculation of ERR default 4 i e 5 p oint scale 0 1 2 3 4
score file	S core ranker's score for each object being ranked has to be used with rank
[]-idv	P rint model performance in test metric on individual ranked lists has to be used w ith test
[]-norm	N ormalize feature vectors similar to norm for training tuning

2. Examples

Go to the [LETOR](#) website and download any of their datasets For instance let's pick MQ2008 from the LETOR 4 0 dataset Suppose we put it under the *RankLib* directory

2.1. Training on held-out data

Type this into the command line

```
> java -jar bin/RankLib.jar -train MQ2008/Fold1/train.txt -test MQ2008/Fold1/test.txt -validate 1
```

What we specified means we want to train a LambdaMART ranker train on the training data and record the model that performs best on the validation data The training metric is NDCG 10 After training is completed evaluate the trained model on the test data in ERR 10 Finally the model will be saved to a file named *mymodel.txt* in the current directory

The parameter `validate` is optional but it often leads to better models. In particular, `validate` is very important for RankNet, MART, LambdaMART, Coordinate Ascent. On the other hand, works pretty well without validation data. As a result, RankLib's implementation of Coordinate Ascent completely ignores the validation data to reduce training time—this is the only exception in RankLib. Starting from **version 2.1 patched**, Coordinate Ascent will utilize validation data if specified just like any other algorithms.

Important Note: `metric2t` (e.g., NDCG, ERR, etc.) only applies to list-wise algorithms. AdaRank, Coordinate Ascent, and LambdaMART, point-wise and pair-wise techniques, MART, RankNet, RankBoost, due to their nature, always use their internal RMSE, pair-wise loss as the optimization criteria. Thus, `metric2t` has no effects on them. ListNet is a special case. Despite being a list-wise algorithm, it has its own optimization criteria as well. Therefore, `metric2t` also has no effect on ListNet.

Tips: -Instead of using a separate validation dataset, you can use `tv`s.

Tips: -Instead of using a separate test dataset, you can use `tt`s.

2.2. k-fold cross validation

Although the LETOR dataset comes with train, test, validation data, in this example, let's pretend that we only had MQ2008 Fold1 train.txt as the only dataset. We now want to do 5-fold cross validation experiment.

a) Sequential partition

```
> java -jar bin/RankLib.jar -train MQ2008/Fold1/train.txt -ranker 4 -kcv 5 -kcvmd models/ -kcvmn
```

The command above will sequentially split the training data into 5 chunks of roughly equal size. The i -th chunk is used as the test data for the i -th fold. The training data for each fold consists of the test data from all other folds.

RankLib will train a Coordinate Ascent model (ranker 4) on each fold that optimizes for NDCG@10. `metric2t` NDCG@10. This model is then evaluated on the corresponding test data for the current fold using ERR@10. `metric2T` ERR@10. After the training process completes, RankLib will report the overall performance across all folds and save all 5 models (one for each fold) to the specified directory `models`: `f1_ca`, `f2_ca`, `f3_ca`, `f4_ca`, and `f5_ca`.

Tips: -You can use `tv`s to reserve a portion of training data in each fold for validation.

b) Randomized partition

Let's say in the training data, ranked lists (e.g., queries) are ordered in a certain way such that sequentially partitioning them might introduce some bias. We want k partitions such that each partition contains a random portion of the input data. We can do this simply by shuffling the order of ranked lists in the training data.

```
java -cp bin/RankLib.jar ciir.umass.edu.features.FeatureManager -input MQ2008/Fold1/train.txt -o
```

The command above will create the shuffled copy that we want called `train.txt.shuffled`, stored in the specified output directory `mydata`. Cross validation can then be done using the command above but with the `shuffled` data instead.

c) How do I obtain the data used in each fold?

To do this, type

```
java -cp bin/RankLib.jar ciir.umass.edu.features.FeatureManager -input MQ2008/Fold1/train.txt.sh
```

This will extract and store the train test data used in each fold which is **exactly the same** as the in memory partitions used for learning. Partitioning is *always* done sequentially.

Tips: Type `java -cp bin/RankLib.jar ciir.umass.edu.features.FeatureManager -for help`

2.3. Evaluating previously trained models

```
> java -jar bin/RankLib.jar -load mymodel.txt -test MQ2008/Fold1/test.txt -metric2T ERR@10
```

This will evaluate the pre trained model stored in *mymodel.txt* on the specified test data using ERR @ 10

2.4. Comparing models

Let's assume our test data *test.txt* contains a set of queries and lists of documents or more precisely their feature vector retrieved for each of the queries using BM25. Let's say we have trained two models: *ca.model.txt* (a Coordinate Ascent model) and *lm.model.txt* (a LambdaMART model) from the same training set. The task is to see if using the Coordinate Ascent model and the LambdaMART model to re-rank these BM25 ranked lists will improve retrieval effectiveness (NDCG @ 10).

It goes like this:

```
> java -jar bin/RankLib.jar -test MQ2008/Fold1/test.txt -metric2T NDCG@10 -idv output/baseline.ndcg.txt
> java -jar bin/RankLib.jar -load ca.model.txt -test MQ2008/Fold1/test.txt -metric2T NDCG@10 -idv output/ca.ndcg.txt
> java -jar bin/RankLib.jar -load lm.model.txt -test MQ2008/Fold1/test.txt -metric2T NDCG@10 -idv output/lm.ndcg.txt
```

Each of the output files specified with *idv* provides the ndcg @ 10 each system achieves for each of the test queries. These files are stored in the *output* directory. Note that these commands are different from the one used in Section 2.3 above which only reports the average measure (e.g. ndcg @ 10) across all queries. These 3 commands, on the other hand, report ndcg @ 10 on each of the queries, not just the average.

Here's an example of an output file showing individual and all query performance levels in terms of the selected metric:

```
NDCG@10 170 0.0
NDCG@10 176 0.6722390270733757
NDCG@10 177 0.4772656487866462
NDCG@10 178 0.539003131276382
NDCG@10 185 0.6131471927654585
NDCG@10 189 1.0
NDCG@10 191 0.6309297535714574
NDCG@10 192 1.0
NDCG@10 194 0.2532778777010656
NDCG@10 197 1.0
NDCG@10 200 0.6131471927654585
NDCG@10 204 0.4772656487866462
NDCG@10 207 0.0
NDCG@10 209 0.123151194370365
NDCG@10 221 0.39038004999210174
NDCG@10 all 0.5193204478059303
```

Now to compare them, do this:

```
> java -cp bin/RankLib.jar ciir.umass.edu.eval.Analyzer -all output/ -base baseline.ndcg.txt > a
```

The output file *analysis.txt* is tab separated. Copy and paste it into any spreadsheet program for easy viewing. Everything should be self explanatory. It looks like this

Overall comparison

System	Performance	Improvement	Win	Loss	p-value
baseline_ndcg.txt	[baseline]	0.093			
LM_ndcg.txt	0.2863	+0.1933 (+207.8%)	9	1	0.03
CA_ndcg.txt	0.5193	+0.4263 (+458.26%)	12	0	0.0

Detailed break down

	[< -100%)	[-100%, -75%)	[-75%, -50%)	[-50%, -25%)	[-25%, 0%)	(0%, +25%]	(+25%, +50%]
LM_ndcg.txt	0	0	1	0	0	4	2
CA_ndcg.txt	0	0	0	0	0	1	6

This output shows performance comparisons of two saved models, CoordinateAscent and LambdaMART, against a baseline. The table shows performance differences and percent improvements between each saved model and the baseline. Also, numbers of queries that were better or worse than baseline and P value for statistical confidence in the better model. Note only queries where performance metrics showed different values are listed in the win/loss columns.

The final part of the output is a simple histogram of performance differences over percent change intervals. Again, only queries that provided differences in metric values are listed and the percent values actually represent difference values between base and test in chosen metric X 100.

Tips: Type `java -cp bin RankLib.jar ciir.umass.edu/eval/Analyzer` for help.

2.5. Using trained models to do re-ranking

Instead of using a model to re-rank documents in some test data and examining the effectiveness of the final rankings, as shown in 2.3, we want to get a hold of the final rankings themselves, i.e., these rankings might serve as input to some other systems.

This can be achieved by

```
> java -jar bin/RankLib.jar -load mymodel.txt -rank MQ2008/Fold1/test.txt -score myscorefile.txt
```

The output file *myscorefile.txt* provides the score that the ranker assigns to each of the documents; higher means more relevant. To obtain the final ranking of documents, simply sort the documents by this score with respect to each query; you will have to write some codes to do this.